


# Meet Frank



HELLO, I'M FRANK! AN INCREMENTAL LEARNING SYSTEM. AND AN ASSISTANT.

I DON'T NEED A TRAINING SET. I LEARN LITTLE BY LITTLE FROM YOUR DECISIONS AND I CAN REPLICATE YOUR BEHAVIOUR. I EVOLVE WITH YOU.

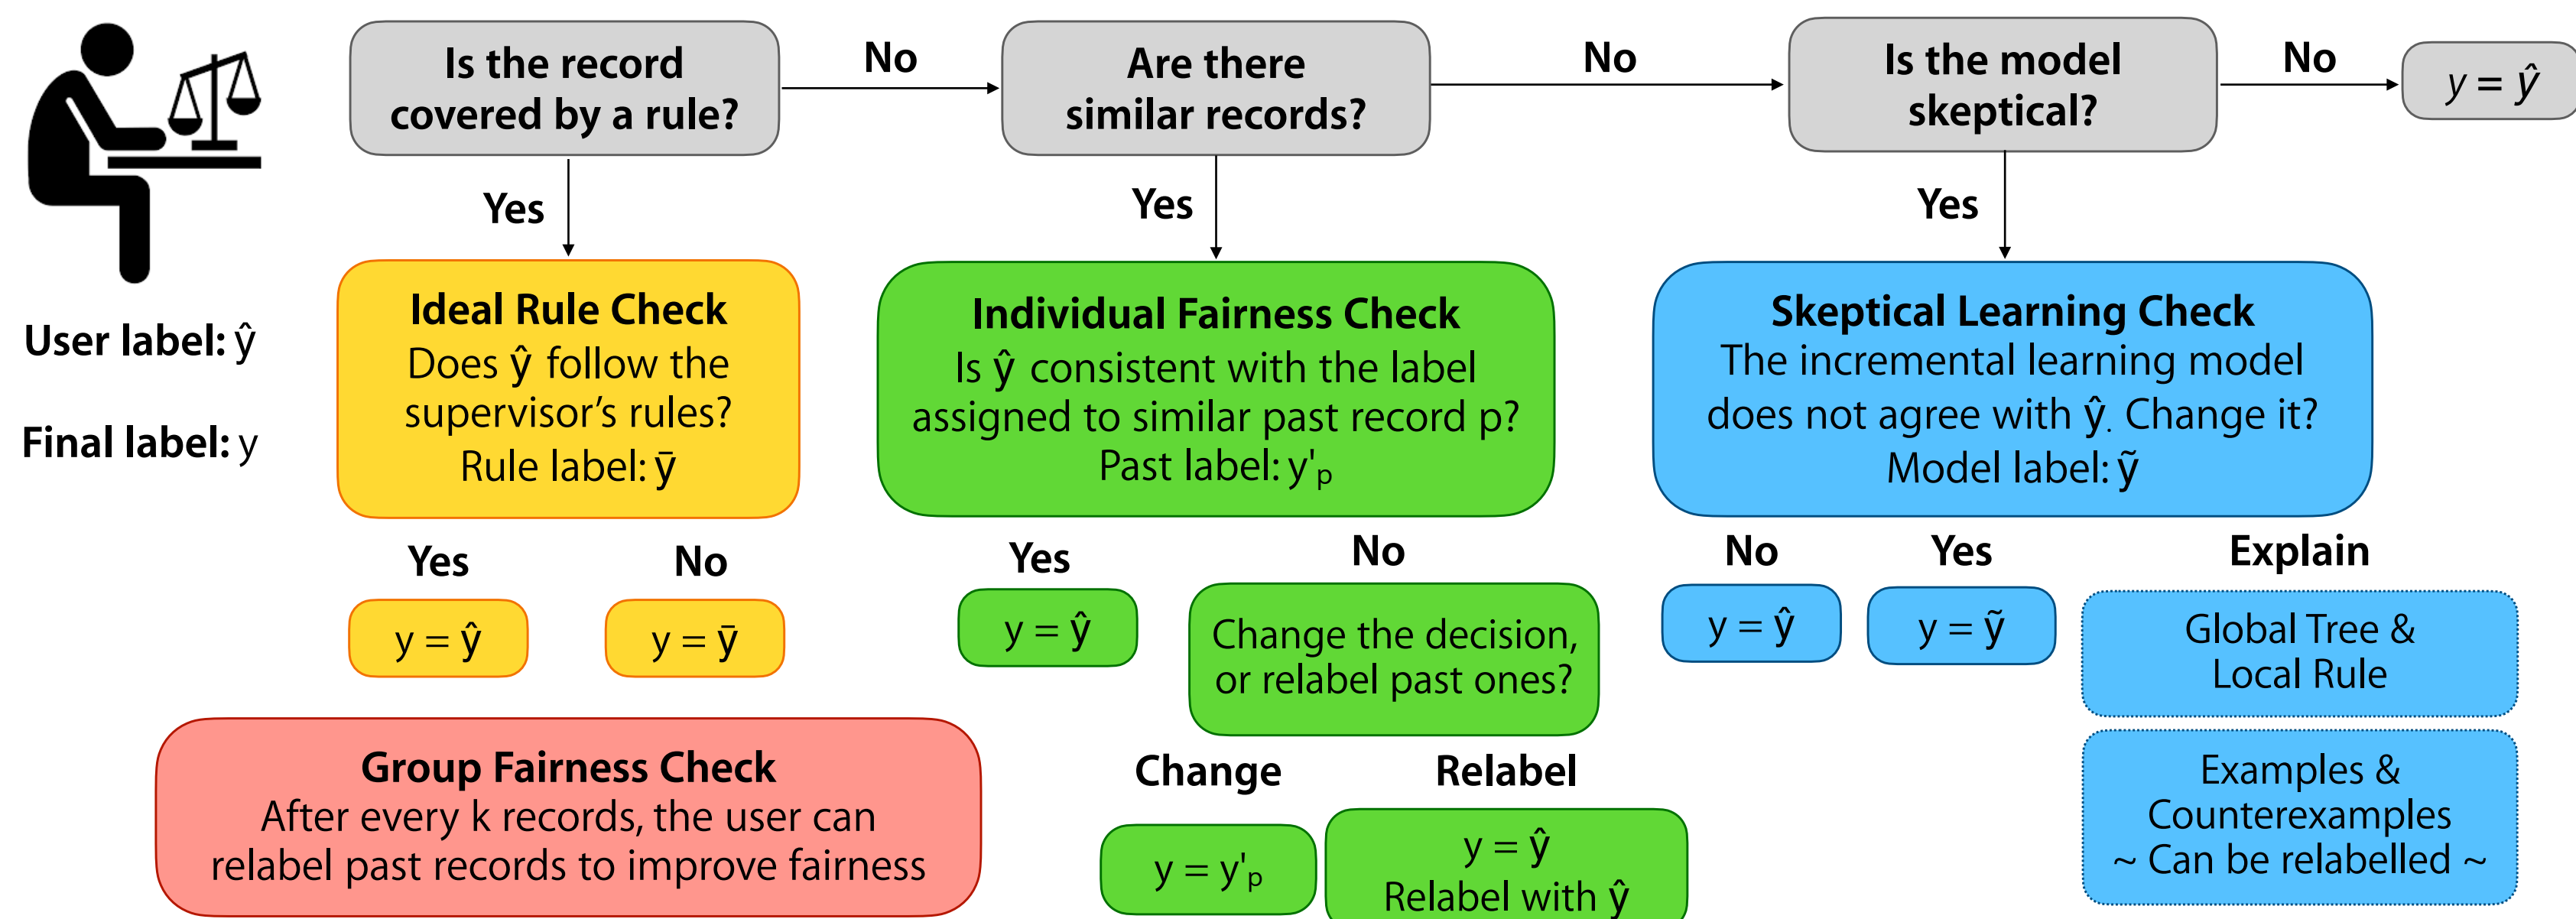
ARE YOU TASKED TO LABEL A DATASET? IF I DON'T AGREE WITH YOUR DECISIONS, I SUGGEST AN ALTERNATIVE TO HELP YOU STAY CONSISTENT.

I CAN ALSO HELP PEOPLE COOPERATE, SIMULATE THEIR FUTURE OR CHALLENGE THEM WITH EDGE-CASES OR INTERESTING SCENARIOS.


I'M FRANK! A FRIEND.

## The now. Frank for Hybrid Decision Making

A user is tasked to label a dataset. Frank helps them stay consistent and fair.



## The future. Applications & Enhancements




### Frank is in charge now!

Frank makes the decisions and, if unsure, asks the user.



### Frank-Human Trust

Frank (dis-)/trusts the user, influencing the behaviour.



### Federated Learning

Frank helps multiple users coordinate, and can find the "weakest link".



### Linguistic Annotation


Frank challenges the two annotators, making them competing on edge-cases.

## The endgame. Going further

### Psychology



Frank helps with psychological diagnosis, evolving the tests following the patient's answers.



### Social Relations

A network of humans is simulated by multiple Franks. How will it evolve?

Last but not least, Frank will also need a good **user interface**.

Frank is a project by Federico Mazzoni.  
Be frank, what do you think? Get in touch!


federico.mazzoni@phd.unipi.it  
<https://github.com/FedericoMz>



UNIVERSITÀ DI PISA



# A Frank System for Co-Evolutionary Hybrid Decision-Making



HELLO, I'M FRANK! AN INCREMENTAL LEARNING SYSTEM. AND AN ASSISTANT.

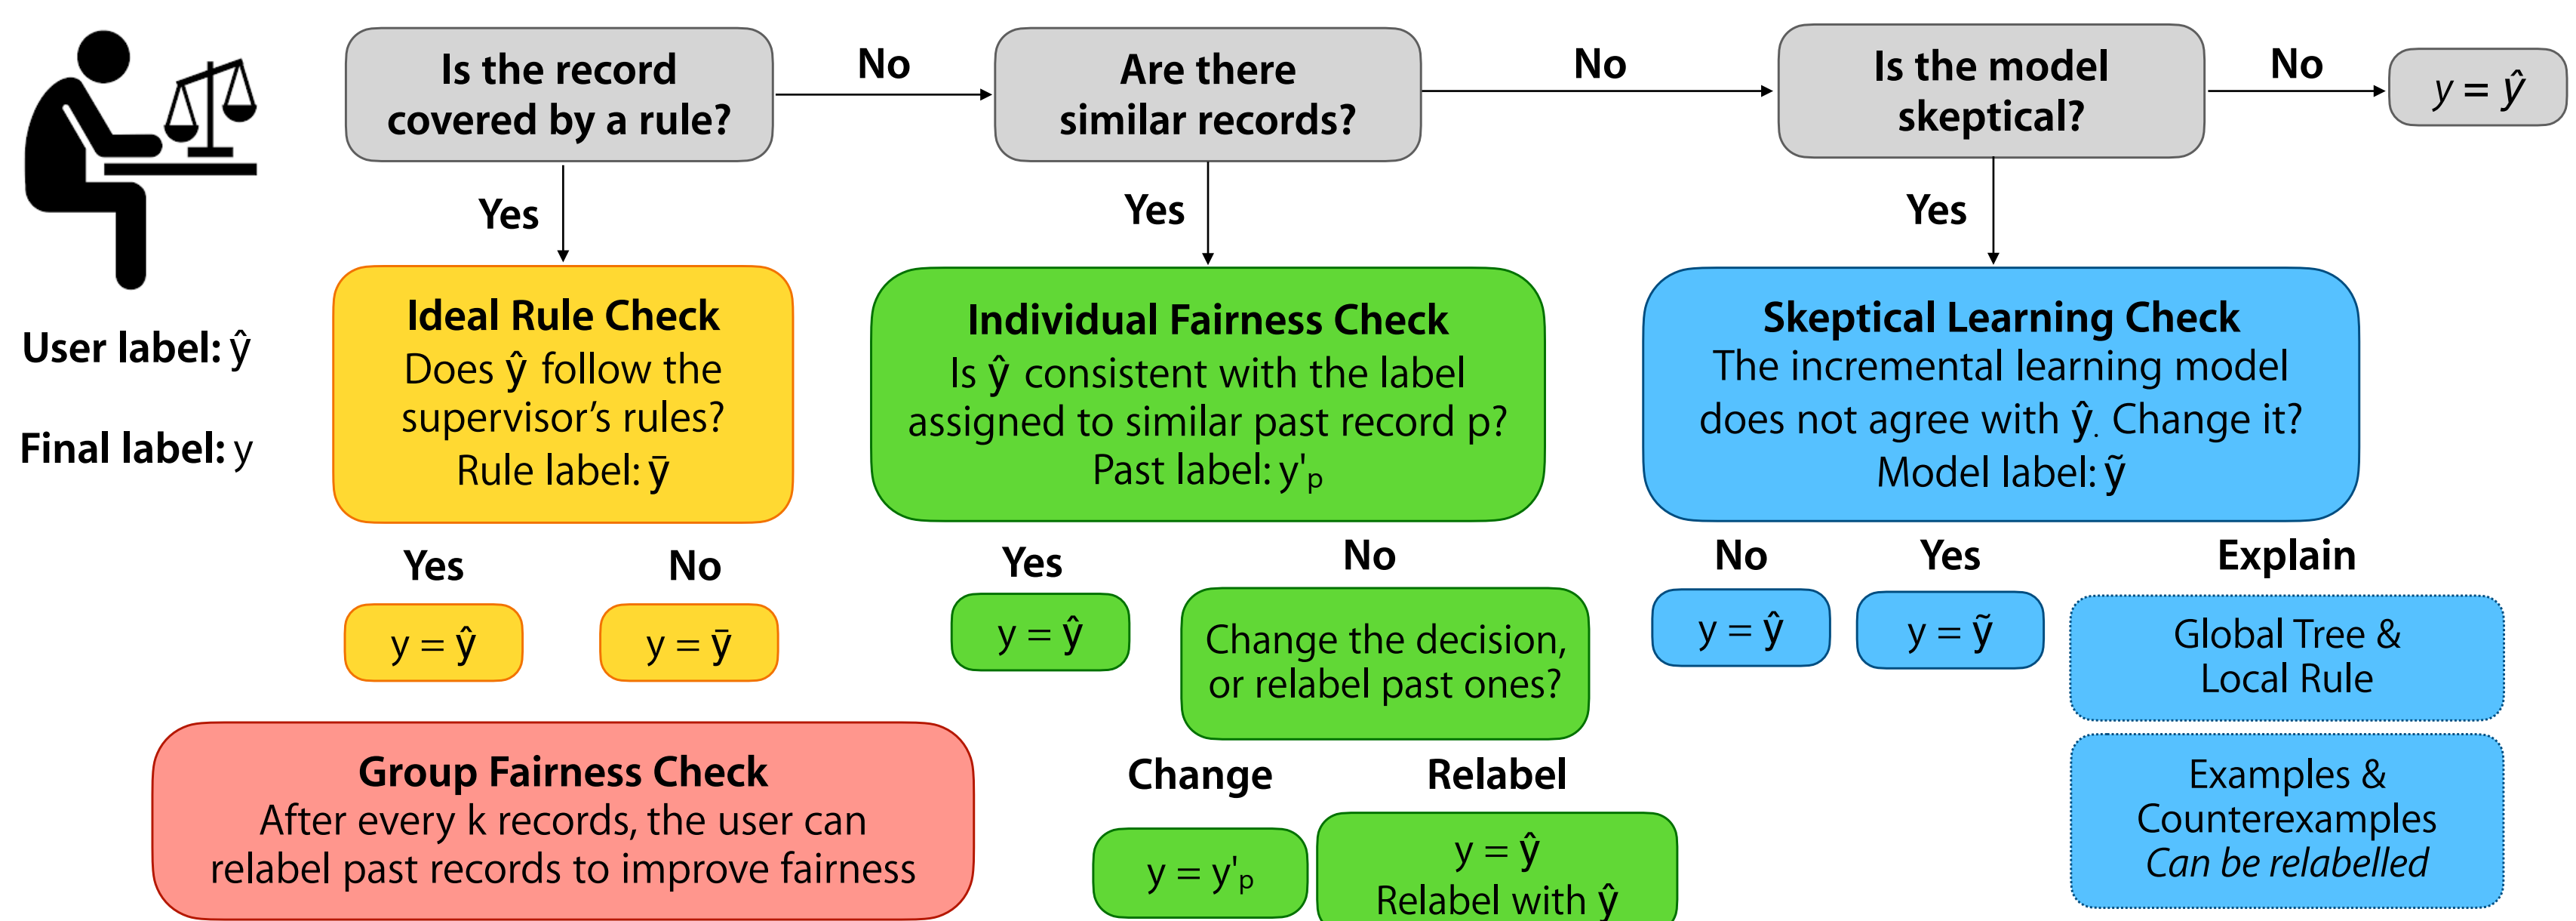
I DON'T NEED A TRAINING SET. I LEARN LITTLE BY LITTLE FROM YOUR DECISIONS AND I CAN REPLICATE YOUR BEHAVIOUR. I EVOLVE WITH YOU.

I'M BASED ON SKEPTICAL LEARNING. ARE YOU TASKED TO LABEL A DATASET? IF I THINK YOU ARE INCONSISTENT WITH YOUR PAST BEHAVIOUR, I CAN SUGGEST AN ALTERNATIVE TO HELP YOU STAY CONSISTENT.

BUT I CAN ALSO HELP YOU FIGHTS YOUR HIDDEN BIASES, AND I CAN PROVIDE MORE AND MORE DETAILED EXPLANATIONS.

I'M FRANK. A FRIEND.

A user is tasked to label a dataset. Frank helps them stay consistent and fair.



**Ideal Rules.** The user most follows decision rules set by the supervisor. The user cannot refuse.

**Individual Fairness.** *Similar individuals should be treated similarly* (e.g., a man and woman with the same background should receive the same label). If they aren't, the user is prompted to change either past decision(s) or the current one. The user cannot refuse.

**Skeptical Learning.** Skepticity is defined following the model's *confidence* (i.e., the prediction probability) and its *certainty* (i.e., its past track record). If the model's skeptical of the user's decision, it asks the user if they want to change it. The user can refuse.

**Group Fairness.** The distribution of the positive label should follow statistical parity between groups (e.g., 20% of men receives a positive label, 20% of women should too). If it isn't, the user is asked if they want to relabel suspicious past record. The user can refuse.

## Data Streams and Incremental Learning

- A. Bifet and R. Gavalda. Adaptive learning from evolving data streams. In IDA, volume 5772 of LNCS, pages 249–260. Springer, 2009.
- P. Domingos et al. Mining high-speed data streams. In KDD, page 71. ACM, 2000.
- M. D. Lange et al. A continual learning survey: Defying forgetting in classification tasks. IEEE Trans. Pattern Anal. Mach. Intell., 44(7):3366–3385, 2022.
- A. Lutz et al. I want to know more—efficient multi-class incremental learning using gaussian processes. Pattern recognition and image analysis, 23:402–407, 2013.
- C. Manapragada, G. I. Webb, and M. Salehi. Extremely fast decision tree. In KDD, pages 1953–1962. ACM, 2018.
- L. Wang, X. Zhang, H. Su, and J. Zhu. A comprehensive survey of continual learning: Theory, method and application. CoRR, abs/2302.00487, 2023.

## Skeptical Learning

- A. Bontempelli, S. Teso, et al. Learning in the wild with incremental skeptical gaussian processes. In IJCAI, pages 2886–2892. ijcai.org, 2020.
- S. Teso, A. Bontempelli, F. Giunchiglia, and A. Passerini. Interactive label cleaning with example-based explanations. In NeurIPS, pages 12966–12977, 2021.
- M. Zeni et al. Fixing mislabeling by human annotators leveraging conflict resolution and prior knowledge. ACM IMWUT, 3(1):32:1–32:23, 2019.
- W. Zhang. Personal context recognition via skeptical learning. In IJCAI, pages 6482–6483. ijcai.org, 2019.

## Hybrid Decision Making

- P. Hemmer, L. Thede, M. V'ossing, J. Jakubik, and N. K'uhl. Learning to defer with limited expert predictions. CoRR, abs/2304.07306, 2023.
- E. Hullermeier and W. Waegeman. Aleatoric and epistemic uncertainty in machine learning: an introduction to concepts and methods. ML, 110(3):457–506, 2021.
- S. Joshi, S. Parbhoo, and F. Doshi-Velez. Pre-emptive learning-to-defer for sequential medical decision-making under uncertainty. CoRR, abs/2109.06312, 2021.
- K. L. Mosier et al. Human decision makers and automated decision aids: Made for each other? In Automation and human performance, pages 201–220. CRC, 2018.
- T. Wang and Q. Lin. Hybrid predictive models: When an interpretable model collaborates with a black-box model. J. Mach. Learn. Res., 22:137:1–137:38, 2021.